Short Keras Introduction

Dr. U. Michelucci (TOELT)

umberto.Michelucci@toelt.ai



TensorFlow



TensorFlow 1.x/2.x

An open source Deep Learning library

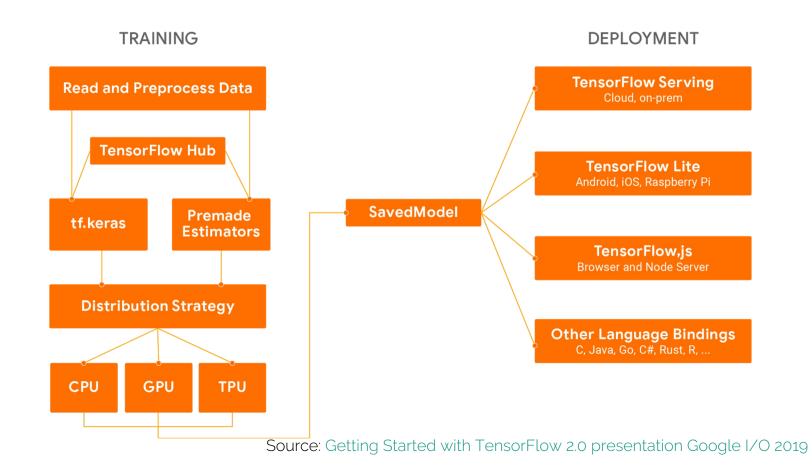
- >1,800 contributors worldwide
- Apache 2.0 license
- Released by Google in 2015

TensorFlow 2.0

- Fasier to learn and use
- For beginners and experts
- Available today

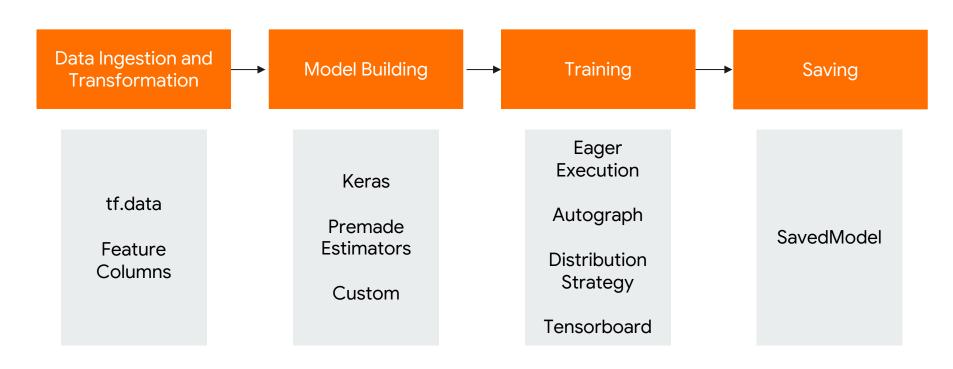


TensorFlow architecture





Training Workflow





TensorFlow 2.x

What is new in TF2.0

- Easy model building with **Keras** and **eager execution** (activated by default in TF2.0).
- Robust model **deployment** in production on any platform.
- Powerful experimentation for research.
- Simplifying the API by **cleaning up deprecated APIs** and reducing duplication (relevant in case you have code developed in TensorFlow 1.X and you need to convert it)
- Load your data using tf.data. Training data is read using input pipelines which are created using tf.data.
- Build, train and validate your model with tf.keras, or use Premade Estimators.
- TensorFlow Hub.
- Run and debug with eager execution, then use tf.function for the benefits of graphs.
- Use Distribution Strategies for distributed training.
- hardware accelerators like CPUs, GPUs, and TPUs; you can enable training workloads to be distributed to single-node/multi-accelerator as well as multi-node/multi-accelerator configurations, including TPU Pods.
- Export to **SavedModel**. TensorFlow will standardize on SavedModel as an interchange format for TensorFlow Serving, TensorFlow Lite, TensorFlow.js, TensorFlow Hub, and more.
- Tensorflow Datasets

tensorflow-gpu version dependencies

Software requirements

The following NVIDIA® software must be installed on your system:

- NVIDIA® GPU drivers ☑ —CUDA 10.0 requires 410.x or higher.
- CUDA® Toolkit ☑ —TensorFlow supports CUDA 10.0 (TensorFlow >= 1.13.0)
- CUPTI ships with the CUDA Toolkit.
- cuDNN SDK <a>□ (>= 7.4.1)
- (Optional) TensorRT 5.0 🖸 to improve latency and throughput for inference on some models.

Check the installed version

import tensorflow as tf
print(tf.__version__)

☆ When you change TF version you need to restart the runtime in Google Colab



TensorFlow 2.0

Usability

- tf.keras as the recommended high-level API.
- Eager execution by default.

```
>>> tf.add(2, 3)
<tf.Tensor: id=2, shape=(), dtype=int32, numpy=5>
```



TensorFlow 2.0

Clarity

- Remove duplicate functionality
- Consistent, intuitive syntax across APIs
- Compatibility throughout the TensorFlow ecosystem



TensorFlow 2.0

Flexibility

- Full lower-level API.
- Internal ops accessible in tf.raw_ops
- Inheritable interfaces for variables, checkpoints, layers.

How to study TF2.0

https://github.com/michelucci/TensorFlow20-Notes/blob/master/howto_study_TF20.md

Basics

- 1. Tensors
- 2. Computational graphs basics (this is relevant when you are going to study how <code>@tf.function</code> works and how it creates a graph in the background. Unless you understand how Computational graphs are working, you will not really understand how <code>@tf.function</code> is working)
- 3. Eager Execution (basics and subtelties)
- 4. tf. Variable (aka no need to initialize them anymore in eager mode)
- 5. Keras basics (Sequential() models, .compile() and .fit() calls)
- 6. Keras functional APIs
- 7. tf.data.Dataset and data processing pipelines

How to study TF2.0

https://github.com/michelucci/TensorFlow20-Notes/blob/master/howto_study_TF20.md

Advanced

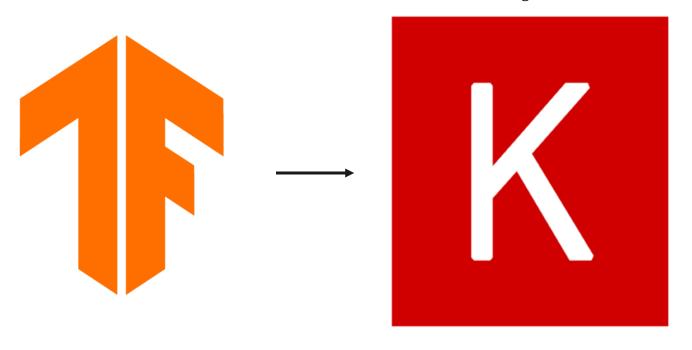
- 1. Model subclassing
- 2. Custom layers in Keras
- 3. Custom callbacks in Keras
- 4. @tf.function
- 5. Keras metrics
- 6. Autodifferentiation (tf.GradientTape()) and custom training loops
- 7. Custom layers and models
- 8. Save and load models
- 9. Using CPUs, GPUs and TPUs (hardware acceleration)

Use a built-in training loop...

```
model.fit(x_train, y_train, epochs=5)
```

Keras

KERAS: high level API



Keras.io (Reference implementation)

import keras

TensorFlow implementation

from tensorflow import keras

★TensorFlow's implementation (a superset, built-in to TF)

For beginners

Model definition

```
model = tf.keras.models.Sequential(...)
```

Model compilation

```
model.compile(...)
```

Model fitting

```
model.fit(...)
```

```
model = tf.keras.models.Sequential([
   tf.keras.layers.Flatten(),
   tf.keras.layers.Dense(512, activation='relu'),
   tf.keras.layers.Dropout(0.2),
   tf.keras.layers.Dense(10, activation='softmax')
])
```

A sequence of layers stacked one after the other

Sequential() model - beginners

```
model = tf.keras.models.Sequential([
  tf.keras.layers.Flatten(),
  tf.keras.layers.Dense(512, activation='relu'),
  tf.keras.layers.Dense(10, activation='softmax')
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test)
```

Sequential() model - experts

```
class MyModel(tf.keras.Model):
    def __init__(self, num_classes=10):
        super(MyModel, self).__init__(name='my_model')
        self.dense_1 = layers.Dense(32, activation='relu')
        self.dense_2 = layers.Dense(num_classes,activation='sigmoid')

def call(self, inputs):
    x = self.dense_1(inputs)
    return self.dense_2(x)
```

The Sequential model is a linear stack of layers.

You can create a **Sequential** model by passing a list of layer instances to the constructor:

```
from keras.models import Sequential
from keras.layers import Dense, Activation

model = Sequential([
    Dense(32, input_shape=(784,)),
    Activation('relu'),
    Dense(10),
    Activation('softmax'),
])
```

You can also simply add layers via the .add() method

```
model = Sequential()
model.add(Dense(32, input_dim=784))
model.add(Activation('relu'))
```

Compilation

Before training a model, you need to configure the learning process

Training

Keras models are trained on Numpy arrays of input data and labels.

model.fit(x_train, y_train, epochs=5)

Layers in Keras

keras.layers

Most used

from keras.layers import Dense, Conv2D,
MaxPooling2D, Dropout, Flatten

keras.layers.Dense
keras.layers.Conv2D
keras.layers.MaxPooling2D
keras.layers.Dropout
keras.layers.Flatten

Example of custom layer

```
class Linear(layers.Layer):
  def init (self, units=32, input dim=32):
    super(Linear, self). init ()
    w init = tf.random normal initializer()
    self.w = tf.Variable(initial_value=w_init(shape=(input_dim, units),
                                              dtype='float32'),
                         trainable=True)
    b init = tf.zeros initializer()
    self.b = tf.Variable(initial_value=b_init(shape=(units,),
                                              dtype='float32'),
                         trainable=True)
  def call(self, inputs):
    return tf.matmul(inputs, self.w) + self.b
x = tf.ones((2, 2))
linear layer = Linear(4, 2)
y = linear layer(x)
```

Custom callback classes

Introduction

A custom callback is a powerful tool to customize the behavior of a Keras model during training, evaluation, or inference, including reading/changing the Keras model.

"In Keras, Callback is a python class meant to be subclassed to provide specific functionality, with a set of methods called at various stages of training (including batch/epoch start and ends), testing, and predicting. "

Custom Callback methods

- on_train_begin: Called at the beginning of training
- on_train_end: Called at the end of training
- on_epoch_begin: Called at the start of an epoch
- on_epoch_end: Called at the end of an epoch
- on_batch_begin: Called right before processing a batch
- on batch end: Called at the end of a batch

Example of Callback

```
class CustomCallback1(keras.callbacks.Callback):
    def on_epoch_end(self, epoch, logs={}):
        print ("Just finished epoch", epoch)
        print (logs)
        return
```

Callback during training

Example of Callback

```
class CustomCallback3(keras.callbacks.Callback):
    def on_epoch_end(self, epoch, logs={}):
        print ("Just finished epoch", epoch)
        print ('Loss evaluated on the validation dataset
        =',logs.get('val_loss'))
        print ('Accuracy reached is', logs.get('acc')) return

>>> Just finished epoch 0

>>> Loss evaluated on the validation dataset = 0.2546206972360611
```

Frequency of logging

```
if (epoch % 10 == 0):
```

Pre-ready callbacks

Classes

```
class BaseLogger: Callback that accumulates epoch averages of metrics.
class CSVLogger: Callback that streams epoch results to a csv file.
class Callback: Abstract base class used to build new callbacks.
class EarlyStopping: Stop training when a monitored quantity has stopped improving.
class History: Callback that records events into a History object.
class LambdaCallback: Callback for creating simple, custom callbacks on-the-fly.
class LearningRateScheduler: Learning rate scheduler.
class ModelCheckpoint: Save the model after every epoch.
class ProgbarLogger: Callback that prints metrics to stdout.
class ReduceLR0nPlateau: Reduce learning rate when a metric has stopped improving.
class RemoteMonitor: Callback used to stream events to a server.
class TensorBoard: Enable visualizations for TensorBoard.
class TerminateOnNaN: Callback that terminates training when a NaN loss is encountered.
```

Multiple Callbacks

Multiple callbacks can be used as a simple list

Keras functional APIs

Keras Functional APIs - an introduction

The Keras functional API is the way to go for defining complex models, such as multi-output models, directed acyclic graphs, models with shared layers or for example residual neural networks.

Example of Keras Functional APIs

Model definition

```
from keras.layers import Input, Dense
from keras.models import Model

inputs = Input(shape=(784,))
output_1 = Dense(64, activation='relu')(inputs)
output_2 = Dense(64, activation='relu')(output_1)
predictions = Dense(10, activation='softmax')(output_2)
model = Model(inputs=inputs, outputs=predictions)
```

Example of Keras Functional APIs

Compiling and training remain the same

Keras Functional APIs - everything is a layer

With the functional API, it is easy to reuse trained models: you can treat any model as if it were a layer, by calling it on a tensor. Note that by calling a model you aren't just reusing the architecture of the model, you are also reusing its weights.

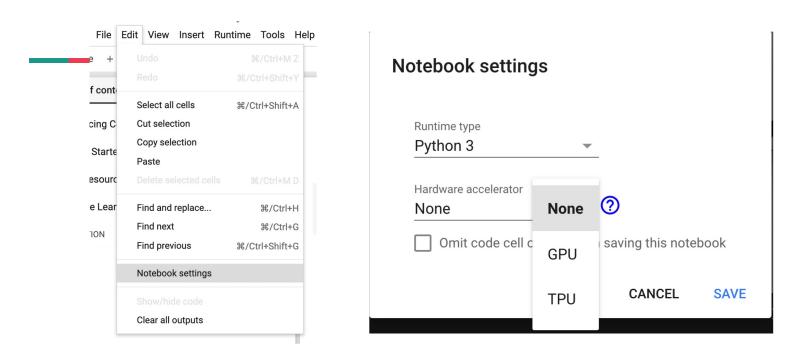
Keras Functional APIs - everything is a layer

This can allow, for instance, to quickly create models that can process sequences of inputs. You could turn an image classification model into a video classification model, in just one line.

```
from keras.layers import TimeDistributed
input_sequences = Input(shape=(20, 784))
processed sequences = TimeDistributed(model)(input_sequences)
```

Hardware acceleration

Using hardware acceleration in Google Colab



Changing the Hardware accelerator setting will make restarting the runtime necessary

Testing presence of a GPU

```
print(tf.test.is_gpu_available())

device_name = tf.test.gpu_device_name()
if device_name != '/device:GPU:0':
    raise SystemError('GPU device not found.')
print('Found GPU at: {}'.format(device_name))
```

Putting operations on a device

```
with tf.device('/gpu:0'):
    # Your code here
with tf.device('/cpu:0'):
    # Your code here
```